

Multiple Things

a bit of everything

Thorsten Beier

February 14, 2019

Outline

Outline

- Deep Learning / Neural Network Frameworks

Outline

- Deep Learning / Neural Network Frameworks
 - Autoencoders

Outline

- Deep Learning / Neural Network Frameworks
 - Autoencoders
 - Variational Autoencoders

Outline

- Deep Learning / Neural Network Frameworks
 - Autoencoders
 - Variational Autoencoders
 - Convolutional Autoencoders

Outline

- Deep Learning / Neural Network Frameworks
 - Autoencoders
 - Variational Autoencoders
 - Convolutional Autoencoders
 - PyTorch

Outline

- Deep Learning / Neural Network Frameworks
 - Autoencoders
 - Variational Autoencoders
 - Convolutional Autoencoders
 - PyTorch
 - Inferno

Outline

- Deep Learning / Neural Network Frameworks
 - Autoencoders
 - Variational Autoencoders
 - Convolutional Autoencoders
 - PyTorch
 - Inferno
- Modern C++

Outline

- Deep Learning / Neural Network Frameworks
 - Autoencoders
 - Variational Autoencoders
 - Convolutional Autoencoders
 - PyTorch
 - Inferno
- Modern C++
 - Why

Outline

- Deep Learning / Neural Network Frameworks
 - Autoencoders
 - Variational Autoencoders
 - Convolutional Autoencoders
 - PyTorch
 - Inferno
- Modern C++
 - Why
 - Improve {Python, R, Julia} with C++

Outline

- Deep Learning / Neural Network Frameworks
 - Autoencoders
 - Variational Autoencoders
 - Convolutional Autoencoders
 - PyTorch
 - Inferno
- Modern C++
 - Why
 - Improve {Python, R, Julia} with C++
 - Rapid Development with C++

Outline

- Deep Learning / Neural Network Frameworks
 - Autoencoders
 - Variational Autoencoders
 - Convolutional Autoencoders
 - PyTorch
 - Inferno
- Modern C++
 - Why
 - Improve {Python, R, Julia} with C++
 - Rapid Development with C++
- Discrete Optimization


Outline

- Deep Learning / Neural Network Frameworks
 - Autoencoders
 - Variational Autoencoders
 - Convolutional Autoencoders
 - PyTorch
 - Inferno
- Modern C++
 - Why
 - Improve {Python, R, Julia} with C++
 - Rapid Development with C++
- Discrete Optimization
 - Energy Minimization


Outline

- Deep Learning / Neural Network Frameworks
 - Autoencoders
 - Variational Autoencoders
 - Convolutional Autoencoders
 - PyTorch
 - Inferno
- Modern C++
 - Why
 - Improve {Python, R, Julia} with C++
 - Rapid Development with C++
- Discrete Optimization
 - Energy Minimization
 - Graphical Models


Autoencoders




Autoencoders




Autoencoders




Autoencoders




Autoencoders




Autoencoders



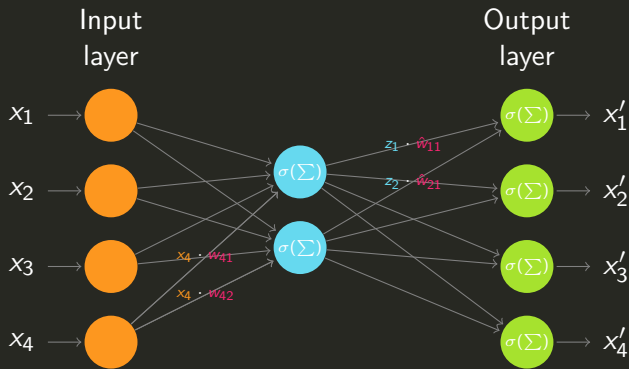
Autoencoders



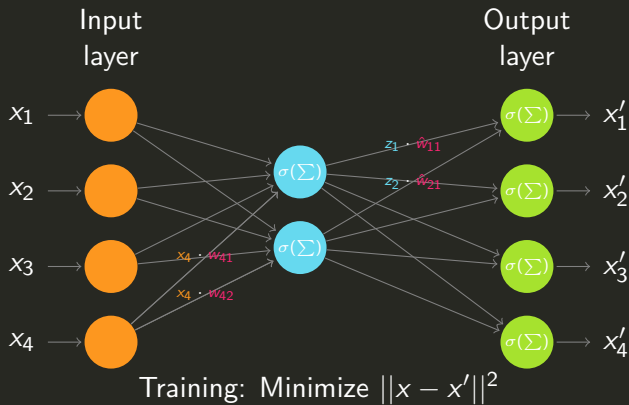
Autoencoders




Autoencoders



Autoencoders




Autoencoders




Test Time: get embeddings for unknown x

Autoencoders



Decode from random embeddings

Convolutional Autoencoder



Deep Learning with PyTorch

```
1  from torch import nn
2  class Autoencoder(nn.Module):
3      def __init__(self):
4          super(Autoencoder, self).__init__()
5
6
7
8
9
10
11
12
13
14
15
16
17
18
19
20
21
22  #
```

Deep Learning with PyTorch

```
1  from pytorch import nn
2  class Autoencoder(nn.Module):
3      def __init__(self):
4          super(autoencoder, self).__init__()
5
6          self.encoder = nn.Sequential(
7              nn.Linear(1024, 256),
8              nn.ReLU(True),
9              nn.Linear(256, 32),
10             nn.ReLU(True))
11
12
13
14
15
16
17
18
19
20
21
22  #
```

Deep Learning with PyTorch

```
1  from pytorch import nn
2  class Autoencoder(nn.Module):
3      def __init__(self):
4          super(autoencoder, self).__init__()
5
6          self.encoder = nn.Sequential(
7              nn.Linear(1024, 256),
8              nn.ReLU(True),
9              nn.Linear(256, 32),
10             nn.ReLU(True))
11
12         self.decoder = nn.Sequential(
13             nn.Linear(32, 256),
14             nn.ReLU(True),
15             nn.Linear(256, 1024),
16             nn.Sigmoid())
17
18
19
20
21
22  #
```

Deep Learning with PyTorch

```
1  from torch import nn
2  class Autoencoder(nn.Module):
3      def __init__(self):
4          super(Autoencoder, self).__init__()
5
6          self.encoder = nn.Sequential(
7              nn.Linear(1024, 256),
8              nn.ReLU(True),
9              nn.Linear(256, 32),
10             nn.ReLU(True))
11
12         self.decoder = nn.Sequential(
13             nn.Linear(32, 256),
14             nn.ReLU(True),
15             nn.Linear(256, 1024),
16             nn.Sigmoid())
17
18     def forward(self, x):
19         x = self.encoder(x)
20
21
22     #
```


Deep Learning with PyTorch

```
1  from pytorch import nn
2  class Autoencoder(nn.Module):
3      def __init__(self):
4          super(autoencoder, self).__init__()
5
6          self.encoder = nn.Sequential(
7              nn.Linear(1024, 256),
8              nn.ReLU(True),
9              nn.Linear(256, 32),
10             nn.ReLU(True))
11
12         self.decoder = nn.Sequential(
13             nn.Linear(32, 256),
14             nn.ReLU(True),
15             nn.Linear(256, 1024),
16             nn.Sigmoid())
17
18     def forward(self, x):
19         x = self.encoder(x)
20         x = self.decoder(x)
21
22     #
```

Deep Learning with PyTorch

```
1  from torch import nn
2  class Autoencoder(nn.Module):
3      def __init__(self):
4          super(Autoencoder, self).__init__()
5
6          self.encoder = nn.Sequential(
7              nn.Linear(1024, 256),
8              nn.ReLU(True),
9              nn.Linear(256, 32),
10             nn.ReLU(True))
11
12         self.decoder = nn.Sequential(
13             nn.Linear(32, 256),
14             nn.ReLU(True),
15             nn.Linear(256, 1024),
16             nn.Sigmoid())
17
18     def forward(self, x):
19         x = self.encoder(x)
20         x = self.decoder(x)
21         return x
22 #
```

```
1 class VariationalAutoencoder(nn.Module):
2     def __init__(self):
3         super(VariationalAutoencoder, self).__init__()
4
5         self.fc1 = nn.Linear(784, 400)
6         self.fc21 = nn.Linear(400, 20)
7         self.fc22 = nn.Linear(400, 20)
8         self.fc3 = nn.Linear(20, 400)
9         self.fc4 = nn.Linear(400, 784)
10
11     def encode(self, x):
12         h1 = F.relu(self.fc1(x))
13         return self.fc21(h1), self.fc22(h1)
14
15     def reparameterize(self, mu, logvar):
16         std = torch.exp(0.5*logvar)
17         eps = torch.randn_like(std)
18         return eps.mul(std).add_(mu)
19
20     def decode(self, z):
21         h3 = F.relu(self.fc3(z))
22         return torch.sigmoid(self.fc4(h3))
23
24     def forward(self, x):
25         mu, logvar = self.encode(x.view(-1, 784))
26         z = self.reparameterize(mu, logvar)
27         return self.decode(z), mu, logvar
```

Deep Learning with PyTorch

```
1 class ConvAutoencoder(nn.Module):
2     def __init__(self):
3         super(ConvAutoencoder, self).__init__()
4         self.encoder = nn.Sequential(
5             nn.Conv2d(1, 16, 3, stride=3, padding=1),
6             nn.ReLU(True),
7             nn.MaxPool2d(2, stride=2),
8             nn.Conv2d(16, 8, 3, stride=2, padding=1),
9             nn.ReLU(True),
10            nn.MaxPool2d(2, stride=1)
11        )
12        self.decoder = nn.Sequential(
13            nn.ConvTranspose2d(8, 16, 3, stride=2),
14            nn.ReLU(True),
15            nn.ConvTranspose2d(16, 8, 5, stride=3, padding=1),
16            nn.ReLU(True),
17            nn.ConvTranspose2d(8, 1, 2, stride=2, padding=1),
18            nn.Tanh()
19        )
20
21    def forward(self, x):
22        x = self.encoder(x)
23        x = self.decoder(x)
24        return x
```

```
1 learning_rate = 1e-4
2 optimizer = torch.optim.Adam(model.parameters(), lr=learning_rate)
3
4
5
6
7
8
9
10
11
12
13
14
15
16
17
18
19
20
21
22
23 #
```

```
1 learning_rate = 1e-4
2 optimizer = torch.optim.Adam(model.parameters(), lr=learning_rate)
3
4 for t in range(500):
5
6     # Forward pass:
7     y_pred = model(x)
8
9
10
11
12
13
14
15
16
17
18
19
20
21
22
23 #
```

```
1 learning_rate = 1e-4
2 optimizer = torch.optim.Adam(model.parameters(), lr=learning_rate)
3
4 for t in range(500):
5
6     # Forward pass:
7     y_pred = model(x)
8
9     # Compute and print loss.
10    loss = loss_fn(y_pred, y)
11    print(t, loss.item())
12
13
14
15
16
17
18
19
20
21
22
23 #
```

```
1 learning_rate = 1e-4
2 optimizer = torch.optim.Adam(model.parameters(), lr=learning_rate)
3
4 for t in range(500):
5
6     # Forward pass:
7     y_pred = model(x)
8
9     # Compute and print loss.
10    loss = loss_fn(y_pred, y)
11    print(t, loss.item())
12
13    # zero all gradients
14    optimizer.zero_grad()
15
16    # Backward pass: compute gradient of
17    # the loss wrt model parameters
18    loss.backward()
19
20
21
22
23 #
```



```
1 learning_rate = 1e-4
2 optimizer = torch.optim.Adam(model.parameters(), lr=learning_rate)
3
4 for t in range(500):
5
6     # Forward pass:
7     y_pred = model(x)
8
9     # Compute and print loss.
10    loss = loss_fn(y_pred, y)
11    print(t, loss.item())
12
13    # zero all gradients
14    optimizer.zero_grad()
15
16    # Backward pass: compute gradient of
17    # the loss wrt model parameters
18    loss.backward()
19
20    # Calling the step function on an
21    # Optimizer makes an update to its parameters
22    optimizer.step()
23    #
```

```
1 learning_rate = 1e-4          # <- fixed learning rate
2 optimizer = torch.optim.Adam(model.parameters(), lr=learning_rate)
3
4 for t in range(500):          # <- to much / to little
5
6     # Forward pass:
7     y_pred = model(x)          # <- no prober data sampling
8
9     # Compute and print loss.
10    loss = loss_fn(y_pred, y)
11    print(t, loss.item())       # <- bad practise for logging
12
13    # zero all gradients
14    optimizer.zero_grad()       # <- boilerplate
15
16    # Backward pass: compute gradient of
17    # the loss wrt model parameters
18    loss.backward()             # <- boilerplate
19
20    # Calling the step function on an
21    # Optimizer makes an update to its parameters
22    optimizer.step()            # <- boilerplate
23    #
```

Inferno

convenience functions/classes around PyTorch

INFERNO



GitHub

<https://github.com/inferno-pytorch/inferno>

Inferno

convenience functions/classes around PyTorch

```
1 from inferno.trainers.basic import Trainer
2 from inferno.trainers.callbacks.logging.tensorboard import TensorboardLogger
3 from inferno.trainers.callbacks.scheduling import AutoLR
4 trainer = Trainer(model)
5 trainer.build_criterion('NLLLoss')
6 trainer.build_metric('CategoricalError')
7 trainer.build_optimizer('Adam')
8 trainer.validate_every((2, 'epochs'))
9 trainer.save_every((5, 'epochs'))
10 trainer.save_to_directory('some/dir')
11 trainer.set_max_num_epochs(10)
12 trainer.register_callback(AutoLR(factor=0.9, patience=(5, 'epochs')))
13 trainer.build_logger(TensorboardLogger(log_scalars_every=(1, 'iteration'),
14                                     log_directory='some/dir'))
15
16 # bind loader
17 trainer.bind_loader('train', train_loader)
18 trainer.bind_loader('validate', validate_loader)
19
20 if USE_CUDA:
21     trainer.cuda()
22
23 trainer.fit()
```

Visualize the training procedure *while training*

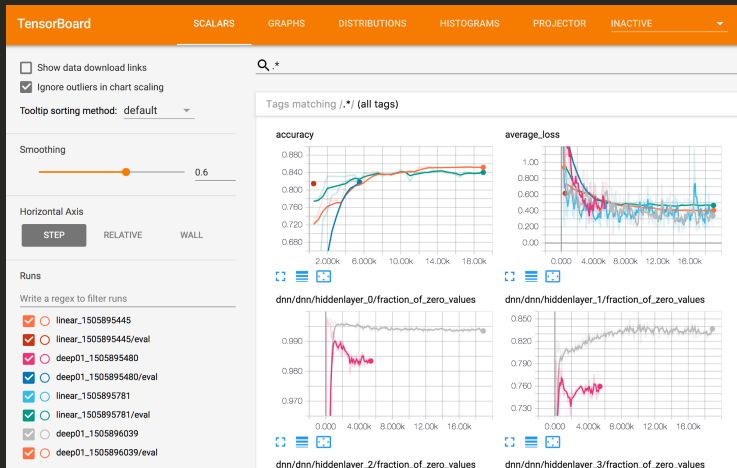


Image Source: <https://towardsdatascience.com/visualizing-your-model-using-tensorboard-796ebb73e98d>

Look at embeddings *while training*

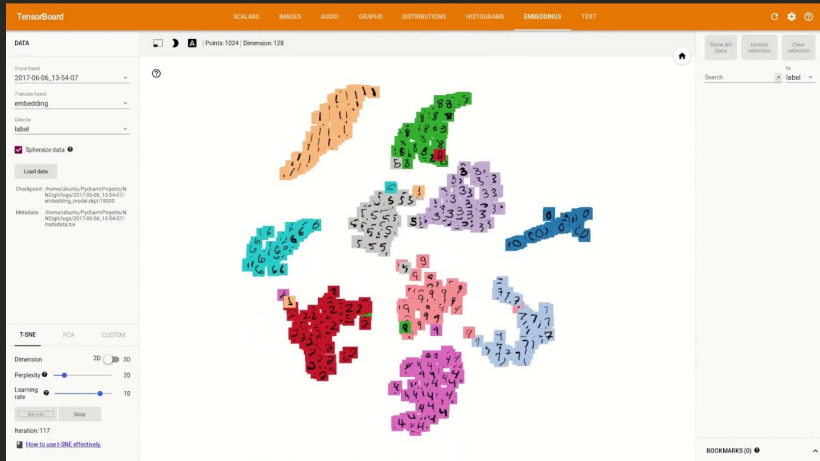


Image Source: <https://www.youtube.com/watch?v=EL0wVFzOgsc>

Why C++



Interpreted Languages are often slow

Interpreted Languages are often slow

```
1 def strange_function(n):
2     s = 0
3     for x in range(n):
4         for y in range(n):
5             for z in range(n):
6                 s += x + 2 * y + z
7     return s
```

Interpreted Languages are often slow


```
1 def strange_function(n):  
2     s = 0  
3     for x in range(n):  
4         for y in range(n):  
5             for z in range(n):  
6                 s += x + 2 * y + z  
7     return s
```


```
1 strange_function(100)    #    0.1 sec  
2 strange_function(1000)  #   100 sec  
3 strange_function(2000)  # >>1000 sec
```

C++ is fast

```
1 auto strange_function(int n)
2 {
3     auto s = 0;
4     for(auto x = 0; x<n; ++x)
5     for(auto y = 0; y<n; ++y)
6     for(auto z = 0; z<n; ++z)
7     {
8         s += x + 2 * y + z;
9     }
10    return s;
11 }
```

```
1 strange_function(100)    # ~0.0 sec
2 strange_function(1000)   #  2.0 sec
3 strange_function(2000)   # 20.0 sec
```





xtensor nd-arrays for C++

Python:

```
np.array([[3, 4], [5, 6]])  
  
arr.reshape([3, 4])  
arr.astype(np.float64)  
  
np.stack([a, b, c], axis=1)  
np.concatenate([a, b, c], axis=1)  
np.squeeze(a)  
np.expand_dims(a, 1)  
np.atleast_3d(a)  
np.split(a, 4, axis=0)  
  
a[:, np.newaxis]  
a[:, 5, 1:]  
a[5:1:-1, :]  
a[..., 3]  
  
np.broadcast(a, [4, 5, 7])  
np.vectorize(f)  
a[a > 5]  
a[[0, 1], [0, 0]]
```

C++:

```
1  xt::xarray<double>({{3, 4}, {5, 6}})  
2  xt::xtensor<double, 2>({{3, 4}, {5, 6}})  
3  arr.reshape({3, 4})  
4  xt::cast<double>(arr)  
5  
6  xt::stack(xt::xtuple(a, b, c), 1)  
7  xt::concatenate(xt::xtuple(a, b, c), 1)  
8  xt::squeeze(a)  
9  xt::expand_dims(a, 1)  
10 xt::atleast_3d(a)  
11 xt::split(a, 4, 0)  
12  
13 xt::view(a, xt::all(), xt::newaxis())  
14 xt::view(a, xt::range(_, 5), xt::range(1, _))  
15 xt::view(a, xt::range(5, 1, -1), xt::all())  
16 xt::strided_view(a, {xt::ellipsis, 3})  
17  
18 xt::broadcast(a, {4, 5, 7})  
19 xt::vectorize(f)  
20 xt::filter(a, a > 5)  
21 xt::index_view(a, {{0, 0}, {1, 0}})
```

C++ Code:

```
1  double sum_of_sines(xt::pyarray<double>& m)
2  {
3      auto sines = xt::sin(m);
4      return std::accumulate(sines.begin(), sines.end(), 0.0);
5  }
6
7  PYBIND11_MODULE(xtensor_python_test, m)
8  {
9      xt::import_numpy();
10     m.def("sum_of_sines", sum_of_sines, "Sum the sines of the input values");
11 }
```

Python Code:

```
1  import numpy as np
2  import xtensor_python_test as xt
3
4  v = np.arange(15).reshape(3, 5)
5  s = xt.sum_of_sines(v)
```

C++ Code:

```
1  using namespace Rcpp;
2
3  // [[Rcpp::plugins(cpp14)]]
4
5  // [[Rcpp::export]]
6  double sum_of_sines(xt::rarray<double>& m)
7  {
8      auto sines = xt::sin(m);
9      return std::accumulate(sines.cbegin(), sines.cend(), 0.0);
10 }
```

R Code:

```
1  v <- matrix(0:14, nrow=3, ncol=5)
2  s <- sum_of_sines(v)
```


C++ Code:

```
1 double sum_of_sines(xt::jarray<double> m)
2 {
3     auto sines = xt::sin(m); // sines does not actually hold values.
4     return std::accumulate(sines.cbegin(), sines.cend(), 0.0);
5 }
6
7 JLCXX_MODULE define_julia_module(jlcxx::Module& mod)
8 {
9     mod.method("sum_of_sines", sum_of_sines);
10 }
```

Julia Code:

```
1 using xtensor_julia_test
2 arr = [[1.0 2.0]
3         [3.0 4.0]]
4 s = sum_of_sines(arr)
```

C++ Notebooks

[JUPYTER](#)[FAQ](#)

Initialize a 2-D array and compute the sum of one of its rows and a 1-D array

```
In [ ]: #include <iostream>

#include "xtensor/xarray.hpp"
#include "xtensor/xio.hpp"
#include "xtensor/xview.hpp"
```

```
In [ ]: xt::xarray<double> arr1
        {{1.0, 2.0, 3.0},
         {2.0, 5.0, 7.0},
         {2.0, 5.0, 7.0}};

        xt::xarray<double> arr2
        {5.0, 6.0, 7.0};

        xt::view(arr1, 1) + arr2
```

```
In [ ]: arr1
```

Initialize a 1-D array and reshape it inplace

```
In [ ]: #include <iostream>
#include "xtensor/xarray.hpp"
#include "xtensor/xio.hpp"
```

```
In [ ]: xt::xarray<int> arr
        {1, 2, 3, 4, 5, 6, 7, 8, 9};

        arr.reshape({3, 3});
```


Try it out: <https://github.com/QuantStack/xtensor>

<https://mybinder.org/v2/gh/QuantStack/xtensor/stable?filepath=notebooks/xtensor.ipynb>

$$f(x) = (x - 0.3)^2 + 1$$

Discrete Optimization


$$f(x) = (x - 0.3)^2 + 1$$



Discrete Optimization

$$f(x) = (x - 0.3)^2 + 1$$


Find $x \in \mathbb{R}$ which minimizes $f(x)$



Discrete Optimization

$$f(x) = (x - 0.3)^2 + 1$$

Find $x \in \mathbb{R}$ which minimizes $f(x)$




$$\hat{x} = \arg \min_{x \in \mathbb{R}} f(x)$$

Discrete Optimization

$$f(x) = (x - 0.3)^2 + 1$$

Find $x \in \mathbb{R}$ which minimizes $f(x)$




$$\hat{x} = \arg \min_{x \in \mathbb{R}} f(x)$$

$$\hat{x} = 0.3$$

Discrete Optimization

$$f(x) = (x - 0.3)^2 + 1$$



Find $x \in \mathbb{R}$ which minimizes $f(x)$


$$\hat{x} = \arg \min_{x \in \mathbb{R}} f(x)$$

$$\hat{x} = 0.3$$

Find integral $x \in \mathbb{N}$ which minimizes $f(x)$

Discrete Optimization

$$f(x) = (x - 0.3)^2 + 1$$



Find $x \in \mathbb{R}$ which minimizes $f(x)$

$$\hat{x} = \arg \min_{x \in \mathbb{R}} f(x)$$


$$\hat{x} = 0.3$$

Find integral $x \in \mathbb{N}$ which minimizes $f(x)$

$$\tilde{x} = \arg \min_{x \in \mathbb{N}} f(x)$$

Discrete Optimization

$$f(x) = (x - 0.3)^2 + 1$$



Find $x \in \mathbb{R}$ which minimizes $f(x)$

$$\hat{x} = \arg \min_{x \in \mathbb{R}} f(x)$$

$$\hat{x} = 0.3$$


Find integral $x \in \mathbb{N}$ which minimizes $f(x)$

$$\tilde{x} = \arg \min_{x \in \mathbb{N}} f(x)$$

$$\tilde{x} = 0$$

Discrete Optimization

$$f(x) = (x - 0.3)^2 + 1$$



Find $x \in \mathbb{R}$ which minimizes $f(x)$

$$\hat{x} = \arg \min_{x \in \mathbb{R}} f(x)$$

$$\hat{x} = 0.3$$

Find integral $x \in \{0, 1, 2, 3, 4\}$
which minimizes $f(x)$

$$\tilde{x} = \arg \min_{x \in \mathbb{N}} f(x)$$

$$\tilde{x} = 0$$

$$x_i \in \{0, 1\}$$

$$E(x_1, x_2, x_3, \dots, x_{N-1}, x_N)$$

$$x_i \in \{0, 1\}$$

$$E(\vec{x})$$

$$x_i \in \{0, 1\}$$

$$E(\vec{x})$$

$$\tilde{\vec{x}} = \arg \min_{\vec{x} \in \{0,1\}^N} E(\vec{x})$$

$$x_i \in \{0, 1\}$$

$$\tilde{\vec{x}} = \arg \min_{\vec{x} \in \{0,1\}^N} E(\vec{x})$$

Enumerating 2^N combinations is impractical

$$x_i \in \{0, 1\}$$

$$E(\vec{x}) = \phi_0(x_0) + \phi_1(x_2) + \phi_1(x_0, x_1) + \phi_2(x_1, x_2) + \phi_3(x_2, x_3)$$

$$\tilde{\vec{x}} = \arg \min_{\vec{x} \in \{0,1\}^N} E(\vec{x})$$

Discrete Optimization

Graphical Models

$$E(\vec{x}) = \phi_1(x_1) + \phi_2(x_2) + \phi_3(x_1, x_2) + \phi_4(x_2, x_3) + \phi_5(x_3, x_4)$$

Discrete Optimization

Graphical Models

$$E(\vec{x}) = \phi_1(x_1) + \phi_2(x_2) + \phi_3(x_1, x_2) + \phi_4(x_2, x_3) + \phi_5(x_3, x_4)$$

Functions as $E(\vec{x})$ are also called *Discrete Graphical Model*

Discrete Optimization

Graphical Models

$$E(\vec{x}) = \phi_1(x_1) + \phi_2(x_2) + \phi_3(x_1, x_2) + \phi_4(x_2, x_3) + \phi_5(x_3, x_4)$$

Functions as $E(\vec{x})$ are also called *Discrete Graphical Model*




Discrete Optimization

Graphical Models

$$E(\vec{x}) = \phi_1(x_1) + \phi_2(x_2) + \phi_3(x_1, x_2) + \phi_4(x_2, x_3) + \phi_5(x_3, x_4)$$

Functions as $E(\vec{x})$ are also called *Discrete Graphical Model*




Discrete Optimization

Graphical Models

$$E(\vec{x}) = \phi_1(x_1) + \phi_2(x_2) + \phi_3(x_1, x_2) + \phi_4(x_2, x_3) + \phi_5(x_3, x_4)$$

Functions as $E(\vec{x})$ are also called *Discrete Graphical Model*




Discrete Optimization

Graphical Models

$$E(\vec{x}) = \phi_1(x_1) + \phi_2(x_2) + \phi_3(x_1, x_2) + \phi_4(x_2, x_3) + \phi_5(x_3, x_4)$$

Functions as $E(\vec{x})$ are also called *Discrete Graphical Model*




- Powerful tool orthogonal to neural networks
- Combinatorial Problems

Discrete Optimization

Graphical Models

$$E(\vec{x}) = \phi_1(x_1) + \phi_2(x_2) + \phi_3(x_1, x_2) + \phi_4(x_2, x_3) + \phi_5(x_3, x_4)$$

Functions as $E(\vec{x})$ are also called *Discrete Graphical Model*



- Powerful tool orthogonal to neural networks
- Combinatorial Problems

Examples:


- Find Optimal seating arrangements for a Table
- Image Segmentation
- Protein Folding

Discrete Optimization

Graphical Models

$$E(\vec{x}) = \phi_1(x_1) + \phi_2(x_2) + \phi_3(x_1, x_2) + \phi_4(x_2, x_3) + \phi_5(x_3, x_4)$$

Functions as $E(\vec{x})$ are also called *Discrete Graphical Model*



- Powerful tool orthogonal to neural networks
- Combinatorial Problems

Examples:

- Find Optimal seating arrangements for a Table
- Image Segmentation
- Protein Folding